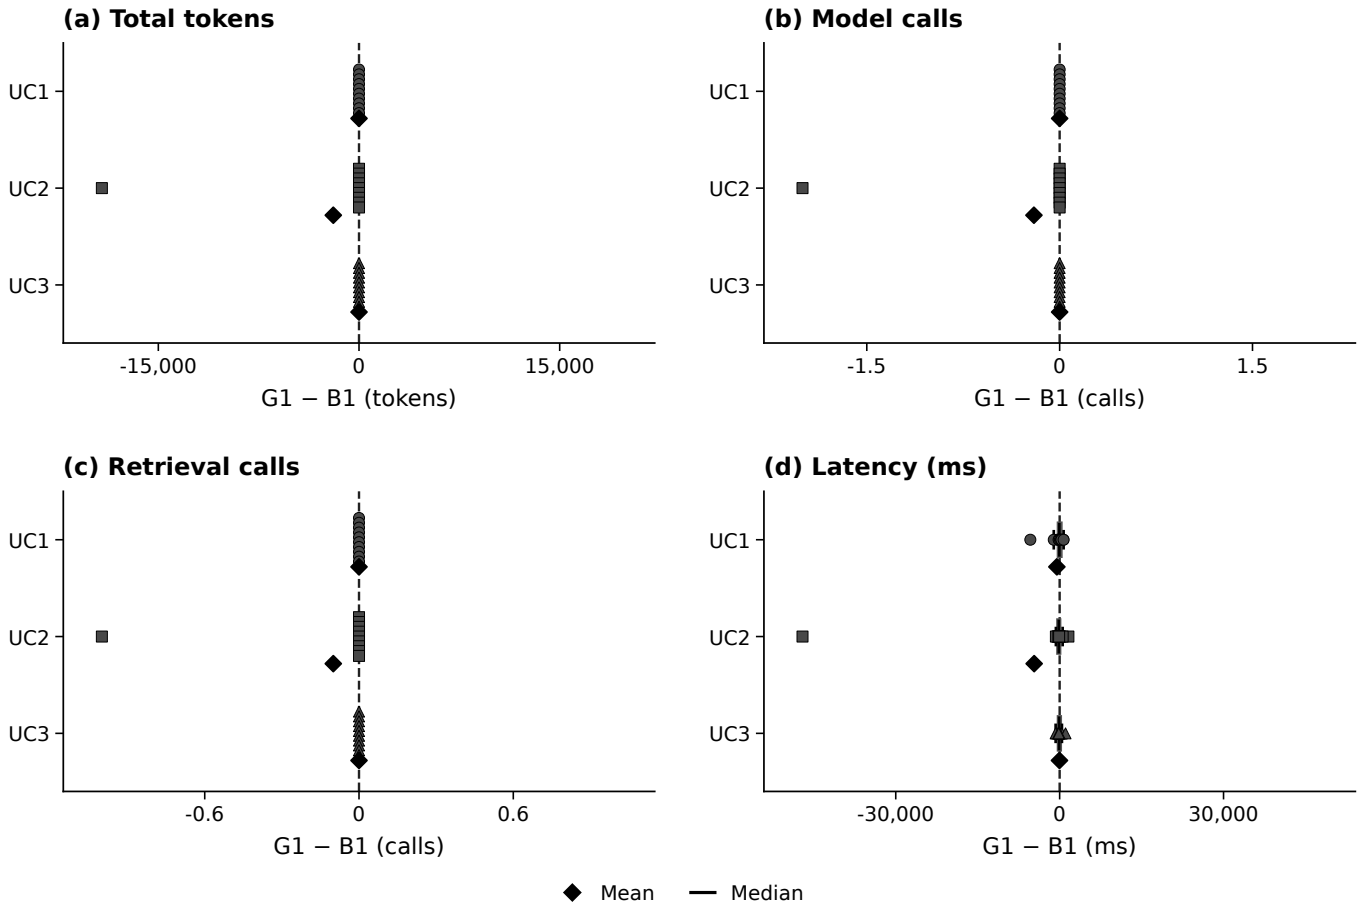
(a) Total tokens
UC1
UC2
UC3
-15,000
0
15,000
G1 − B1 (tokens)
(b) Model calls
UC1
UC2
UC3
-1.5
0
1.5
G1 − B1 (calls)
(c) Retrieval calls
UC1
UC2
UC3
-0.6
0
0.6
G1 − B1 (calls)
(d) Latency (ms)
UC1
UC2
UC3
-30,000
0
30,000
G1 − B1 (ms)
Mean
Median